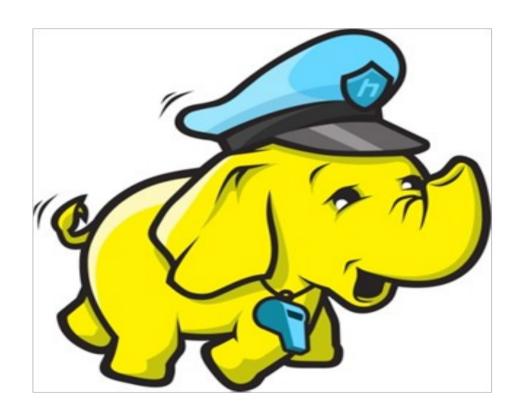
HDFS & MR Administration



For online Hadoop training, send mail to neeraj.ymca.2k6@gmail.com

Agenda

Namenode in safemode Check the health of HDFS Setting quota on no. of files Clearing quota on no. of files Setting quota on space Clearing quota on space **Balance your HDFS** Deal with corrupt blocks Move the corrupt files to lost+found Delete the corrupt files Monitoring MR Jobs Kill MR jobs from command terminal

Namenode in safemode

After start-up, Namenode remains in safemode for some time.

You can't do write operation, while Namenode is in safemode.

```
🔞 🖹 🗈 Terminal
neeraj@ubuntu:~/local_cluster_home/hadoop-1.2.1/bin$ pwd
/home/neeraj/local_cluster_home/hadoop-1.2.1/bin
neeraj@ubuntu:~/local_cluster_home/hadoop-1.2.1/bin$ ./hadoop dfsadmin -safemode get
Safe mode is OFF
neeraj@ubuntu:~/local_cluster_home/hadoop-1.2.1/bin$ ./hadoop dfsadmin -safemode enter
Safe mode is ON
neeraj@ubuntu:~/local_cluster_home/hadoop-1.2.1/bin$ ./hadoop dfsadmin -safemode get
Safe mode is ON
neeraj@ubuntu:~/local_cluster_home/hadoop-1.2.1/bin$ ./hadoop dfsadmin -safemode leave
Safe mode is OFF
neeraj@ubuntu:~/local_cluster_home/hadoop-1.2.1/bin$ ./hadoop dfsadmin -safemode get
Safe mode is OFF
neeraj@ubuntu:~/local_cluster_home/hadoop-1.2.1/bin$
```

Check the health of HDFS

./hadoop fsck /

```
风 🖨 📵 🏻 Terminal
neeraj@ubuntu:~/local cluster home/hadoop-1.2.1/bin$ pwd
/home/neeraj/local cluster home/hadoop-1.2.1/bin
neeraj@ubuntu:~/local cluster home/hadoop-1.2.1/bin$ ./hadoop fsck /
FSCK started by neeraj from /127.0.0.1 for path / at Sun Aug 18 14:50:51 IST 2013
..Status: HEALTHY
 Total size: 127 B
Total dirs: 11
Total files: 2
 Total blocks (validated): 2 (avg. block size 63 B)
Minimally replicated blocks: 2 (100.0 %)
Over-replicated blocks:
                               0 (0.0 %)
Under-replicated blocks:
                               0 (0.0 %)
Mis-replicated blocks:
                               0 (0.0 %)
Default replication factor:
Average block replication:
                               1.0
 Corrupt blocks:
                               0 (0.0 %)
Missing replicas:
Number of data-nodes:
Number of racks:
FSCK ended at Sun Aug 18 14:50:51 IST 2013 in 3 milliseconds
The filesystem under path '/' is HEALTHY
neeraj@ubuntu:~/local cluster home/hadoop-1.2.1/bin$
```

Check the health of HDFS

./hadoop fsck /dir1

```
风 🖨 📵 🏻 Terminal
neeraj@ubuntu:~/local cluster home/hadoop-1.2.1/bin$ pwd
/home/neeraj/local cluster home/hadoop-1.2.1/bin
neeraj@ubuntu:~/local cluster home/hadoop-1.2.1/bin$ ./hadoop fsck /dir1
FSCK started by neeraj from /127.0.0.1 for path /dir1 at Sun Aug 18 15:46:19 IST 2013
.Status: HEALTHY
Total size: 123 B
Total dirs:
Total files:
Total blocks (validated): 1 (avg. block size 123 B)
Minimally replicated blocks: 1 (100.0 %)
Over-replicated blocks:
                            0 (0.0 %)
Under-replicated blocks: 0 (0.0 %)
                          0 (0.0 %)
Mis-replicated blocks:
Default replication factor:
Average block replication: 1.0
Corrupt blocks:
Missing replicas:
                              0 (0.0 %)
Number of data-nodes:
Number of racks:
FSCK ended at Sun Aug 18 15:46:19 IST 2013 in 10 milliseconds
The filesystem under path '/dir1' is HEALTHY
neeraj@ubuntu:~/local cluster home/hadoop-1.2.1/bin$
```

Check the block health of HDFS

./hadoop fsck /dir1 -files -blocks

```
风 🖨 📵 Terminal
neeraj@ubuntu:~/local cluster home/hadoop-1.2.1/bin$ pwd
/home/neeraj/local cluster home/hadoop-1.2.1/bin
neeraj@ubuntu:~/local cluster home/hadoop-1.2.1/bin$ ./hadoop fsck /dir1 -files -blocks
FSCK started by neeraj from /127.0.0.1 for path /dir1 at Sun Aug 18 15:06:23 IST 2013
/dir1 <dir>
/dir1/sample hadoop.txt 123 bytes, 1 block(s): OK
0. blk -8754226978439846901 1011 len=123 repl=1
Status: HEALTHY
 Total size:
              123 B
 Total dirs:
 Total files: 1
Total blocks (validated):
                               1 (avg. block size 123 B)
Minimally replicated blocks:
                               1 (100.0 %)
 Over-replicated blocks:
                               0 (0.0 %)
Under-replicated blocks:
                               0 (0.0 %)
Mis-replicated blocks:
                               0 (0.0 %)
Default replication factor:
 Average block replication:
                               1.0
 Corrupt blocks:
Missing replicas:
                               0 (0.0 %)
Number of data-nodes:
Number of racks:
FSCK ended at Sun Aug 18 15:06:23 IST 2013 in 1 milliseconds
The filesystem under path '/dir1' is HEALTHY
neeraj@ubuntu:~/local cluster home/hadoop-1.2.1/bin$
```

Setting quota on no. of files

Many small files, increase the load of Namenode.

So we can put a restriction on maximum no. of files in a directory.

Clearing quota on no. of files

If we have put the restriction on number of files in a directory.

We can remove this restriction, if required.

Setting quota on space

We can put restriction on the size of a directory in HDFS

So that DFS Admin can control the usage of storage.

We can use m for MB, g for GB, t for TB

Clearing quota on space

If we have put the restriction on storage for directory.

We can remove this restriction, if required.

Balance your HDFS

Whenever we add/remove nodes from Hadoop cluster.

We need to balance the Hadoop cluster.

./hadoop balancer

```
🔞 🖨 📵 🏻 Terminal
neeraj@ubuntu:~/local_cluster_home/hadoop-1.2.1/bin$ pwd
/home/neeraj/local_cluster_home/hadoop-1.2.1/bin
neeraj@ubuntu:~/local_cluster_home/hadoop-1.2.1/bin$ ./hadoop balancer
                       Iteration# Bytes Already Moved Bytes Left To Move Bytes Being Moved
Time Stamp
13/08/18 15:10:23 INFO net.NetworkTopology: Adding a new node: /default-rack/127.0.0.1:50010
13/08/18 15:10:23 INFO balancer.Balancer: 0 over utilized nodes:
13/08/18 15:10:23 INFO balancer.Balancer: 1 under utilized nodes: 127.0.0.1:50010
The cluster is balanced. Exiting...
Balancing took 1.019 seconds
neeraj@ubuntu:~/local_cluster_home/hadoop-1.2.1/bin$
```

Deal with corrupt blocks

Blocks may get corrupt due to hardware /software failures.

MR jobs may get failed due to corrupt input data.

We need to either move or delete the corrupted data.

Corrupt blocks can be used for root cause analysis.

After deletion of corrupted blocks, the replication factor for some files decreases.

HDFS creates new replica to balance the replication factor.

Move the corrupt files to lost+found

./hadoop fsck -move

```
    □    □    Terminal

neeraj@ubuntu:~/local cluster home/hadoop-1.2.1/bin$ pwd
/home/neeraj/local cluster home/hadoop-1.2.1/bin
neeraj@ubuntu:~/local_cluster_home/hadoop-1.2.1/bin$ ./hadoop_fsck -move
FSCK started by neeraj from /127.0.0.1 for path / at Sun Aug 18 15:15:56 IST 2013
..Status: HEALTHY
 Total size: 127 B
 Total dirs: 12
 Total files:
 Total blocks (validated): 2 (avg. block size 63 B)
 Minimally replicated blocks: 2 (100.0 %)
 Over-replicated blocks:
                               0 (0.0 %)
 Under-replicated blocks:
                               0 (0.0 %)
 Mis-replicated blocks:
                               0 (0.0 %)
 Default replication factor:
 Average block replication:
                               1.0
 Corrupt blocks:
 Missing replicas:
                               0 (0.0 %)
 Number of data-nodes:
Number of racks:
FSCK ended at Sun Aug 18 15:15:56 IST 2013 in 34 milliseconds
The filesystem under path '/' is HEALTHY
neeraj@ubuntu:~/local_cluster_home/hadoop-1.2.1/bin$
```

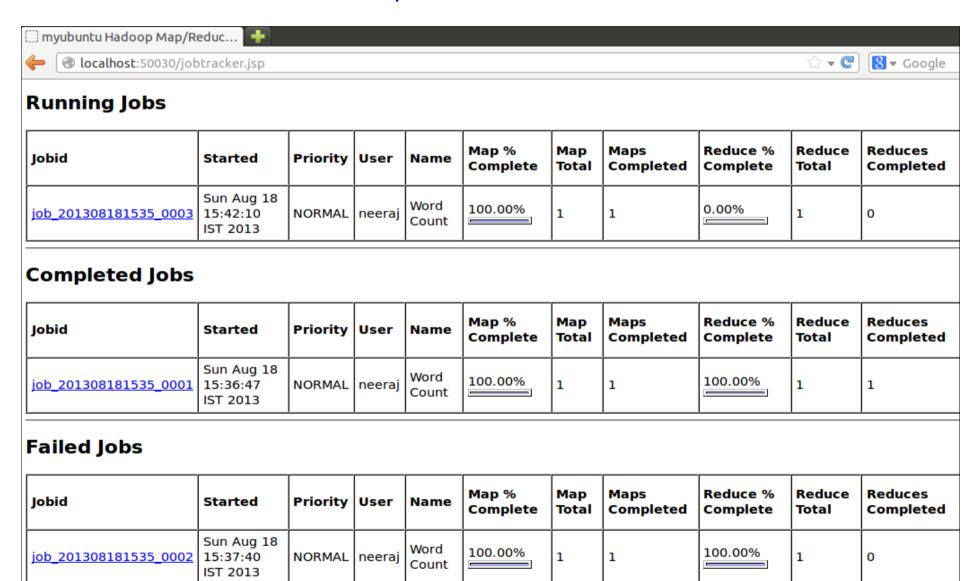
Delete the corrupt blocks

./hadoop fsck -delete

```
🙉 🖨 📵 Terminal
neeraj@ubuntu:~/local cluster home/hadoop-1.2.1/bin$ pwd
/home/neeraj/local cluster home/hadoop-1.2.1/bin
neeraj@ubuntu:~/local cluster home/hadoop-1.2.1/bin$ ./hadoop fsck -delete
FSCK started by neeraj from /127.0.0.1 for path / at Sun Aug 18 15:18:45 IST 2013
..Status: HEALTHY
 Total size: 127 B
 Total dirs: 12
 Total files:
 Total blocks (validated): 2 (avg. block size 63 B)
 Minimally replicated blocks: 2 (100.0 %)
 Over-replicated blocks:
                               0 (0.0 %)
Under-replicated blocks:
                               0 (0.0 %)
 Mis-replicated blocks:
                               0 (0.0 %)
 Default replication factor:
 Average block replication:
                               1.0
 Corrupt blocks:
Missing replicas:
                               0 (0.0 %)
 Number of data-nodes:
Number of racks:
FSCK ended at Sun Aug 18 15:18:45 IST 2013 in 4 milliseconds
The filesystem under path '/' is HEALTHY
neeraj@ubuntu:~/local_cluster_home/hadoop-1.2.1/bin$
```

Monitoring MR Jobs

http://localhost:50030

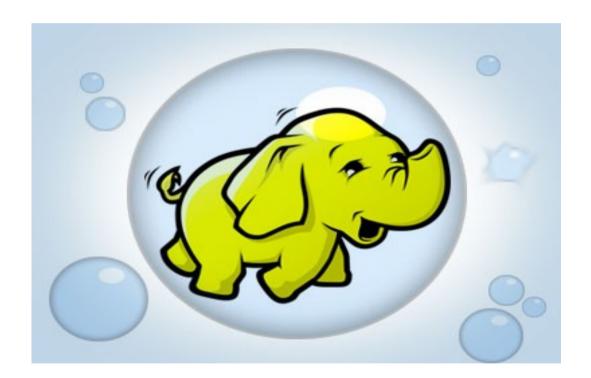


Kill MR jobs from command terminal

./hadoop job -kill < job Id >

```
neeraj@ubuntu:~/local_cluster_home/hadoop-1.2.1/bin$ pwd
/home/neeraj/local_cluster_home/hadoop-1.2.1/bin$ pwd
/home/neeraj/local_cluster_home/hadoop-1.2.1/bin$ ./hadoop jar ~/local_cluster_home/job_jars/WordCountJar.jar
examples.WordCount /dir1/sample_hadoop.txt /wordcount_op2
13/08/18 15:37:38 WARN mapred.JobClient: Use GenericOptionsParser for parsing the arguments. Applications should
implement Tool for the same.
13/08/18 15:37:39 INFO input.FileInputFormat: Total input paths to process: 1
13/08/18 15:37:39 INFO util.NativeCodeLoader: Loaded the native-hadoop library
13/08/18 15:37:39 WARN snappy.LoadSnappy: Snappy native library not loaded
13/08/18 15:37:40 INFO mapred.JobClient: Running job: job_201308181535_0002
13/08/18 15:37:41 INFO mapred.JobClient: map 0% reduce 0%
```

...Thanks...



For online Hadoop training, send mail to neeraj.ymca.2k6@gmail.com